



CSE120

Principles of Operating Systems

Prof Yuanyuan (YY) Zhou

Lecture 4: Threads



Announcement

- Project 0 Due
- Project 1 out
- Homework 1 due on Thursday
 - Submit it to Gradescope online

Processes

- Recall that a process includes many things
 - An address space (defining all the code and data pages)
 - OS resources (e.g., open files) and accounting information
 - Execution state (PC, SP, regs, etc.)
- Creating a new process is costly because of all of the data structures that must be allocated and initialized
 - Recall struct proc in Solaris
- Communicating between processes is costly because most communication goes through the OS
 - Overhead of system calls and copying data

Parallel Programs

- To execute these programs we need to
 - Create several processes that execute in parallel
 - Have the OS schedule these processes in parallel (logically or physically)
- This situation is **very inefficient**
 - **Space**: PCB, page tables, etc.
 - **Time**: create data structures, fork and copy addr space, etc.
- Solution: possible to have cooperating “*processes*”?

Rethinking Processes

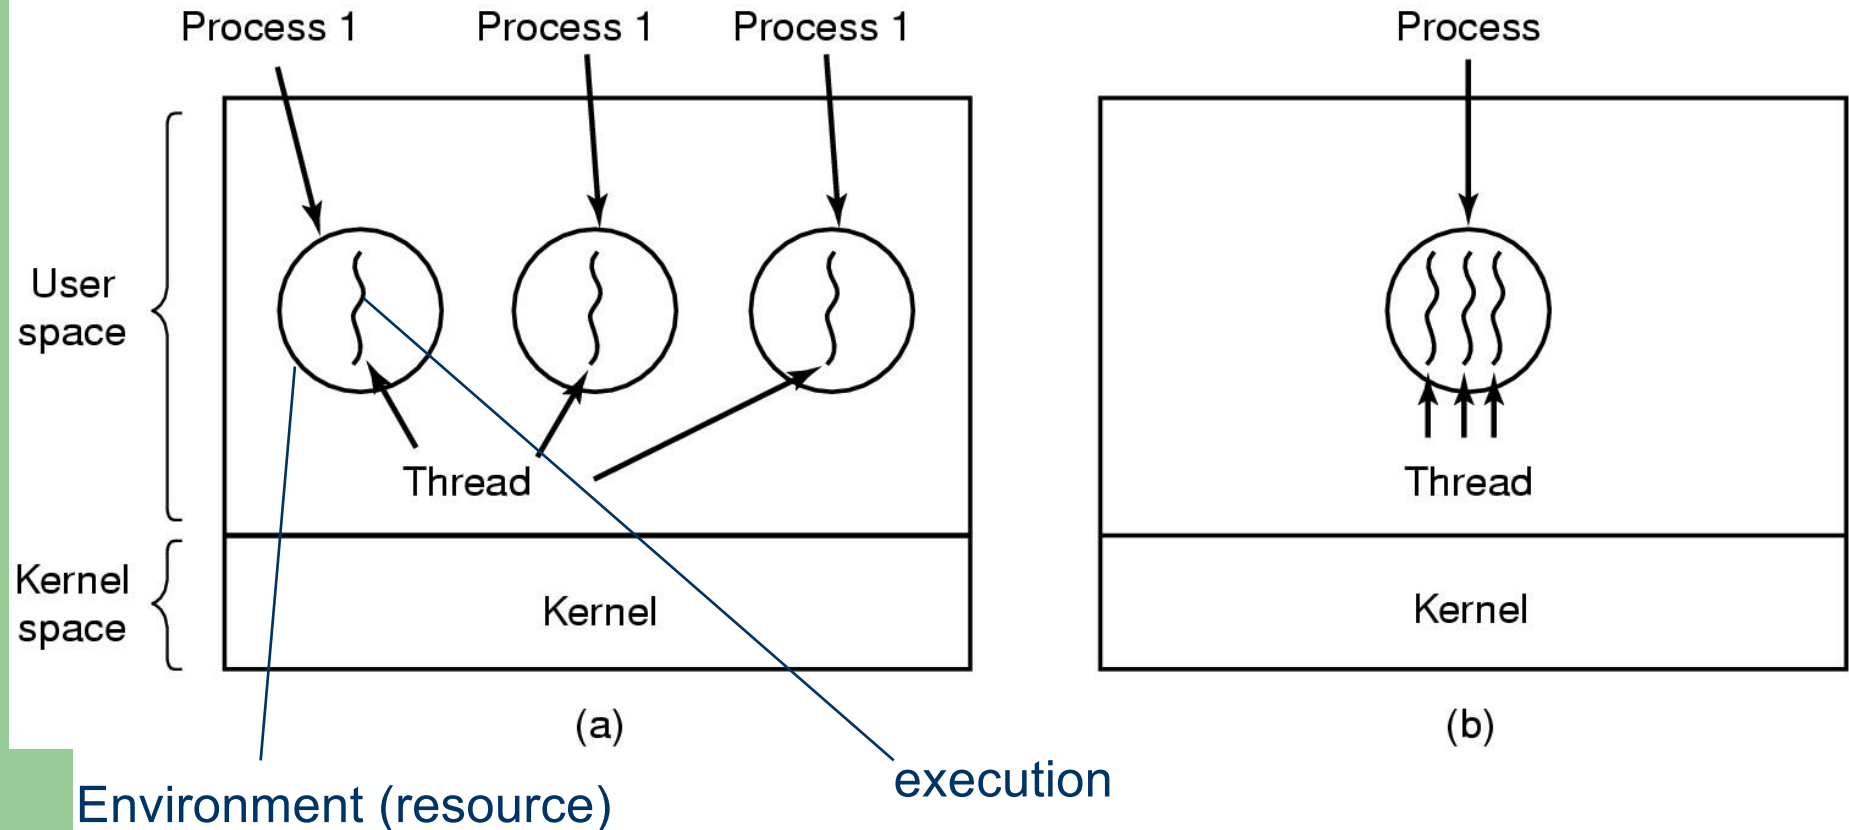
- **Key idea:** Why don't we separate the concept of a process from its execution state?
 - **Process:** address space, privileges, resources, etc.
 - **Execution state:** PC, SP, registers
- Exec state also called **thread of control**, or **thread**

Threads

- Thread vs Process
 - A **thread** defines a sequential execution stream within a process (PC, SP, registers)
 - A **process** defines the address space and general process attributes (everything but threads of execution)
- A thread is bound to a single process
 - A process, however, can have multiple threads
- Threads become the unit of scheduling
 - Processes are now the **containers** in which threads execute

Threads: Lightweight Processes

A sequential execution stream within a process

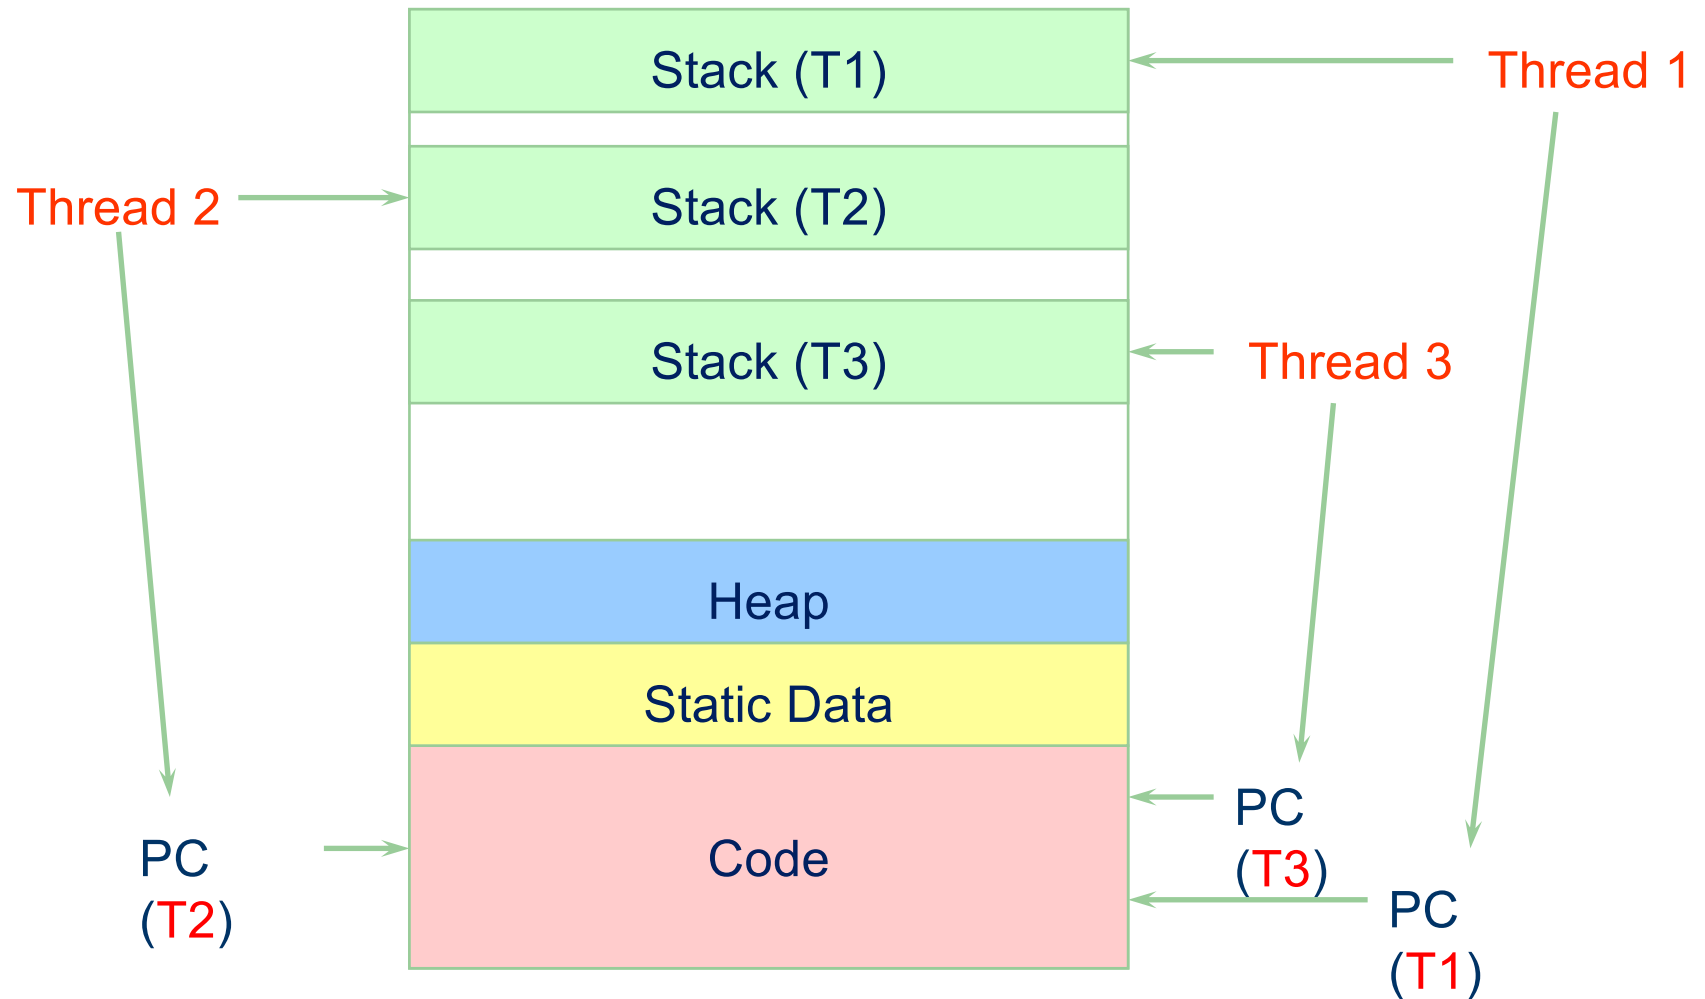


- (a) Three processes each with one thread
- (b) One process with three threads

The Thread Model

- Shared information
 - Processor info: parent process, time, etc
 - Memory: global data, heap, page table, and stats, etc
 - I/O and file: communication ports, directories and file descriptors, etc
- Private state
 - State (ready, running and blocked)
 - Registers
 - Program counter
 - Execution stack
 - **Why?**
- Each thread execute separately

Threads in a Process



Analogy

- Process: 3 projects for different classes (CSE120, CSE140, CSE110)
 - Each one has different text book, different web pages, different TAs/Instructors
- Threads: 3 activities in CSE120 (Homework, Lectures, Projects)
 - Share the same concepts (textbook)
 - Share TA/Tutors
 - All of them are going on in parallel (within one quarter)
 - Each has their own things, too



A 2-min Explainer Video

- <https://www.youtube.com/watch?v=Dhf-DYO1K78>

Threads: Concurrent Servers

- Using fork() to create new processes to handle requests in parallel is overkill for such a simple task
- Recall our forking Web server:

```
while (1) {  
    int sock = accept();  
    if ((child_pid = fork()) == 0) {  
        Handle client request  
        Close socket and exit  
    } else {  
        Close socket  
    }  
}
```

Threads: Concurrent Servers

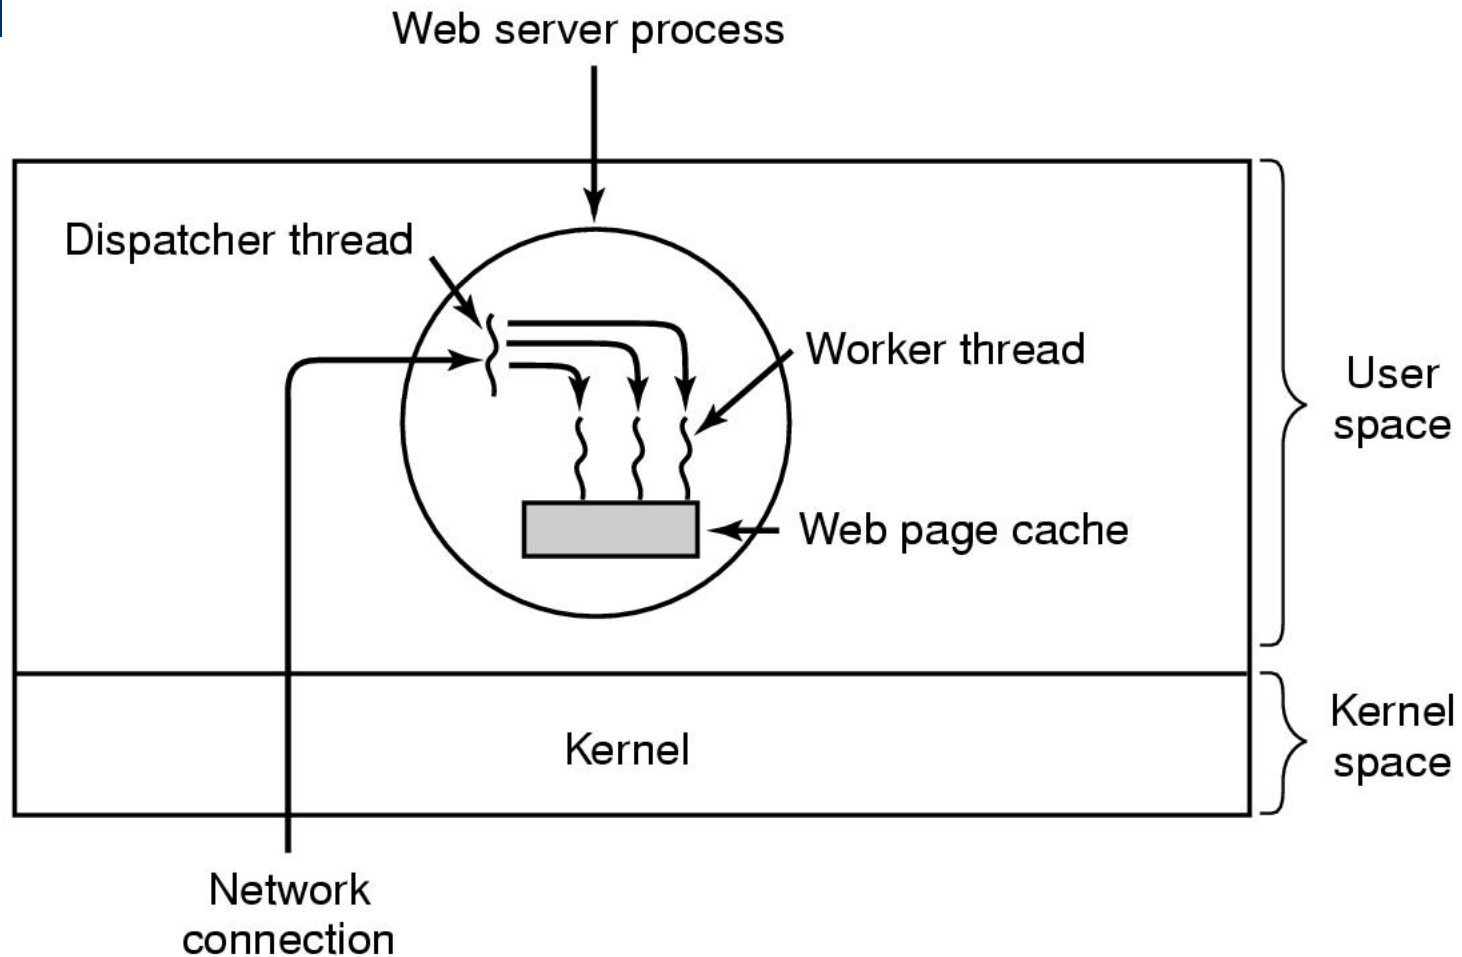
- Instead, we can create a new thread for each request

```
web_server() {  
    while (1) {  
        int sock = accept();  
        thread_fork(handle_request, sock);  
    }  
}
```


Difference from fork()?

```
handle_request(int sock) {  
    Process request  
    close(sock);  
}
```

Thread Usage: Web Server

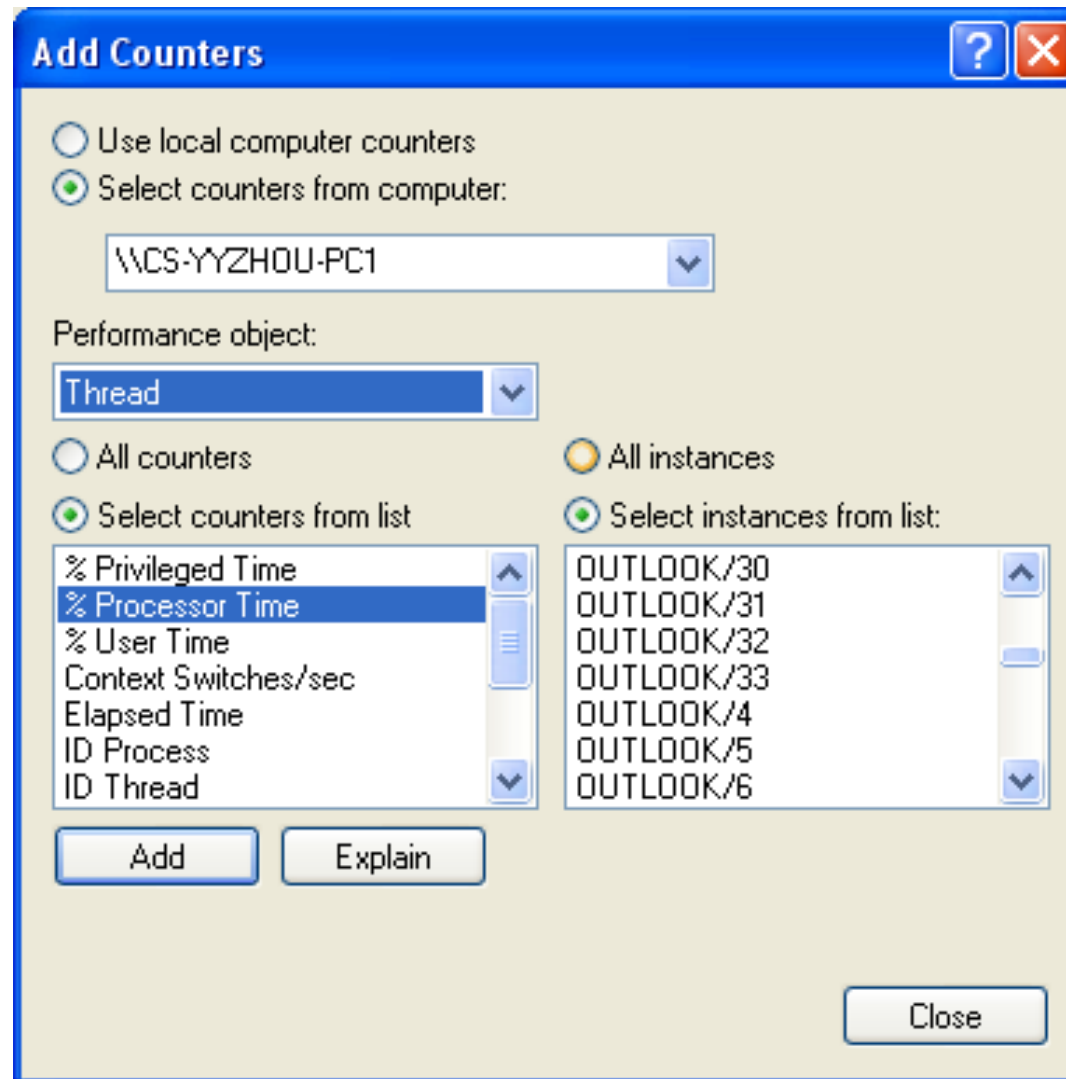


Thread Usage: word processor

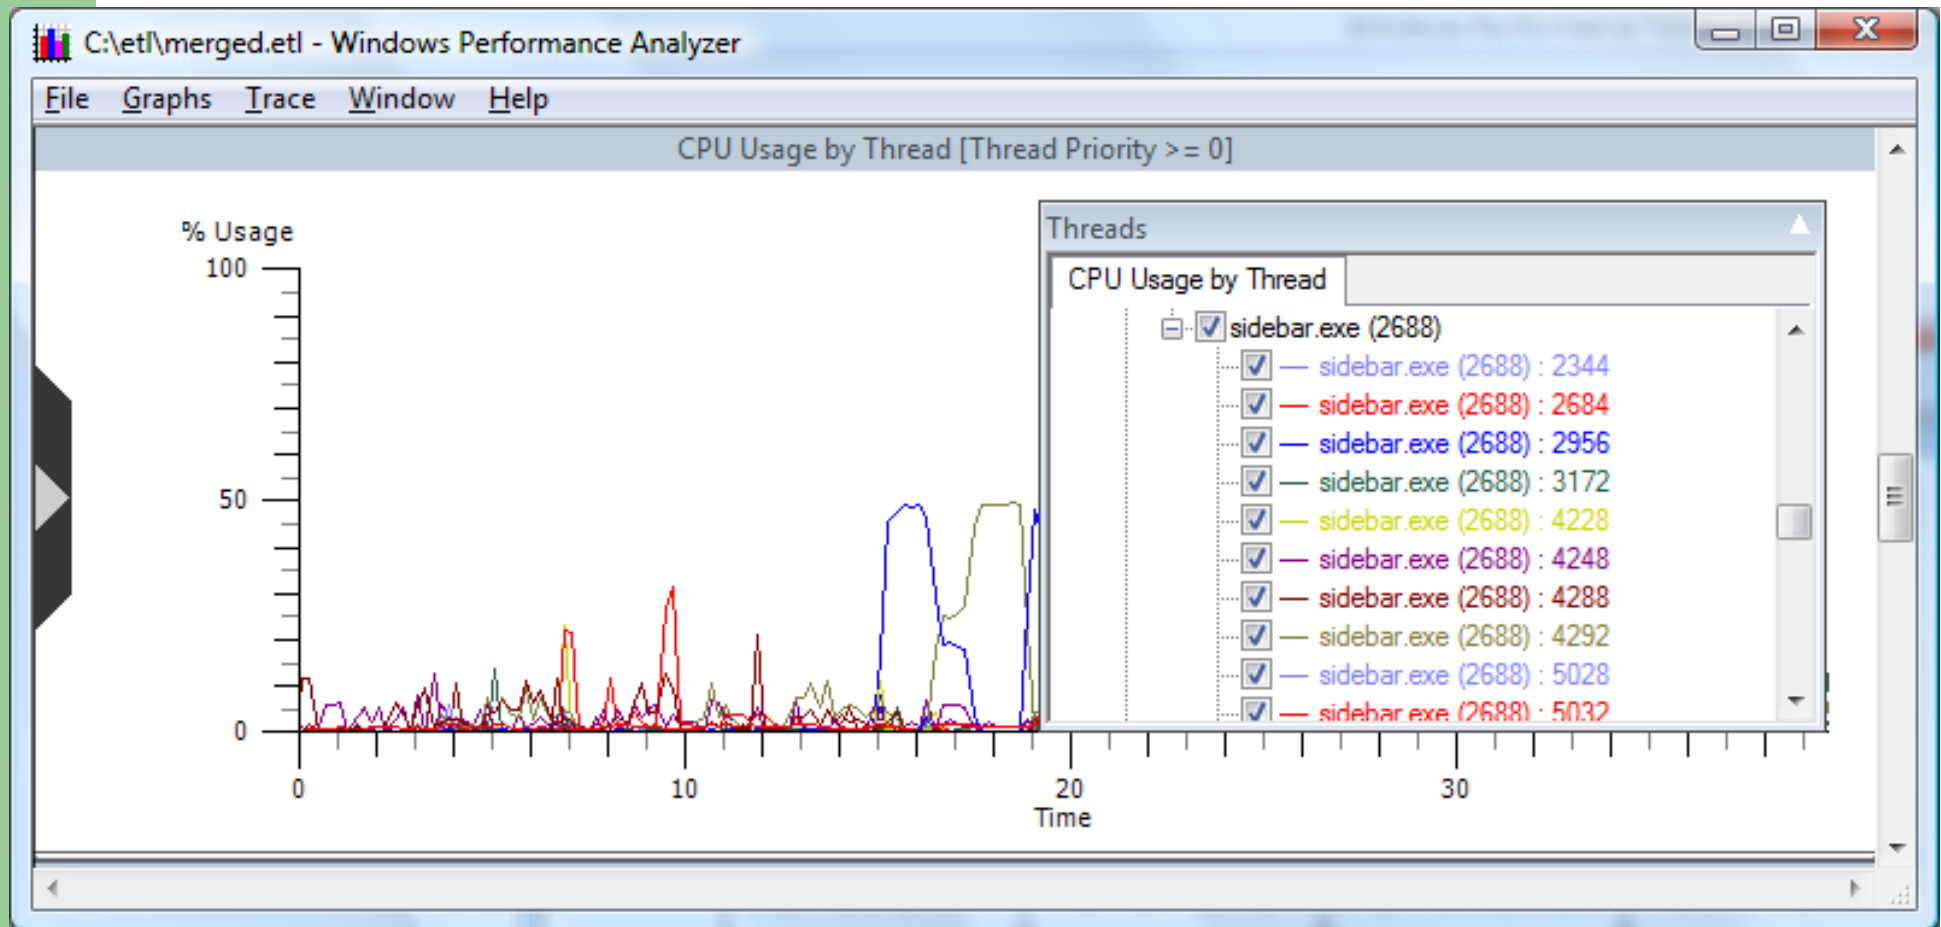


- A thread can wait for I/O, while the other threads can still running.
- What if it is single-threaded?

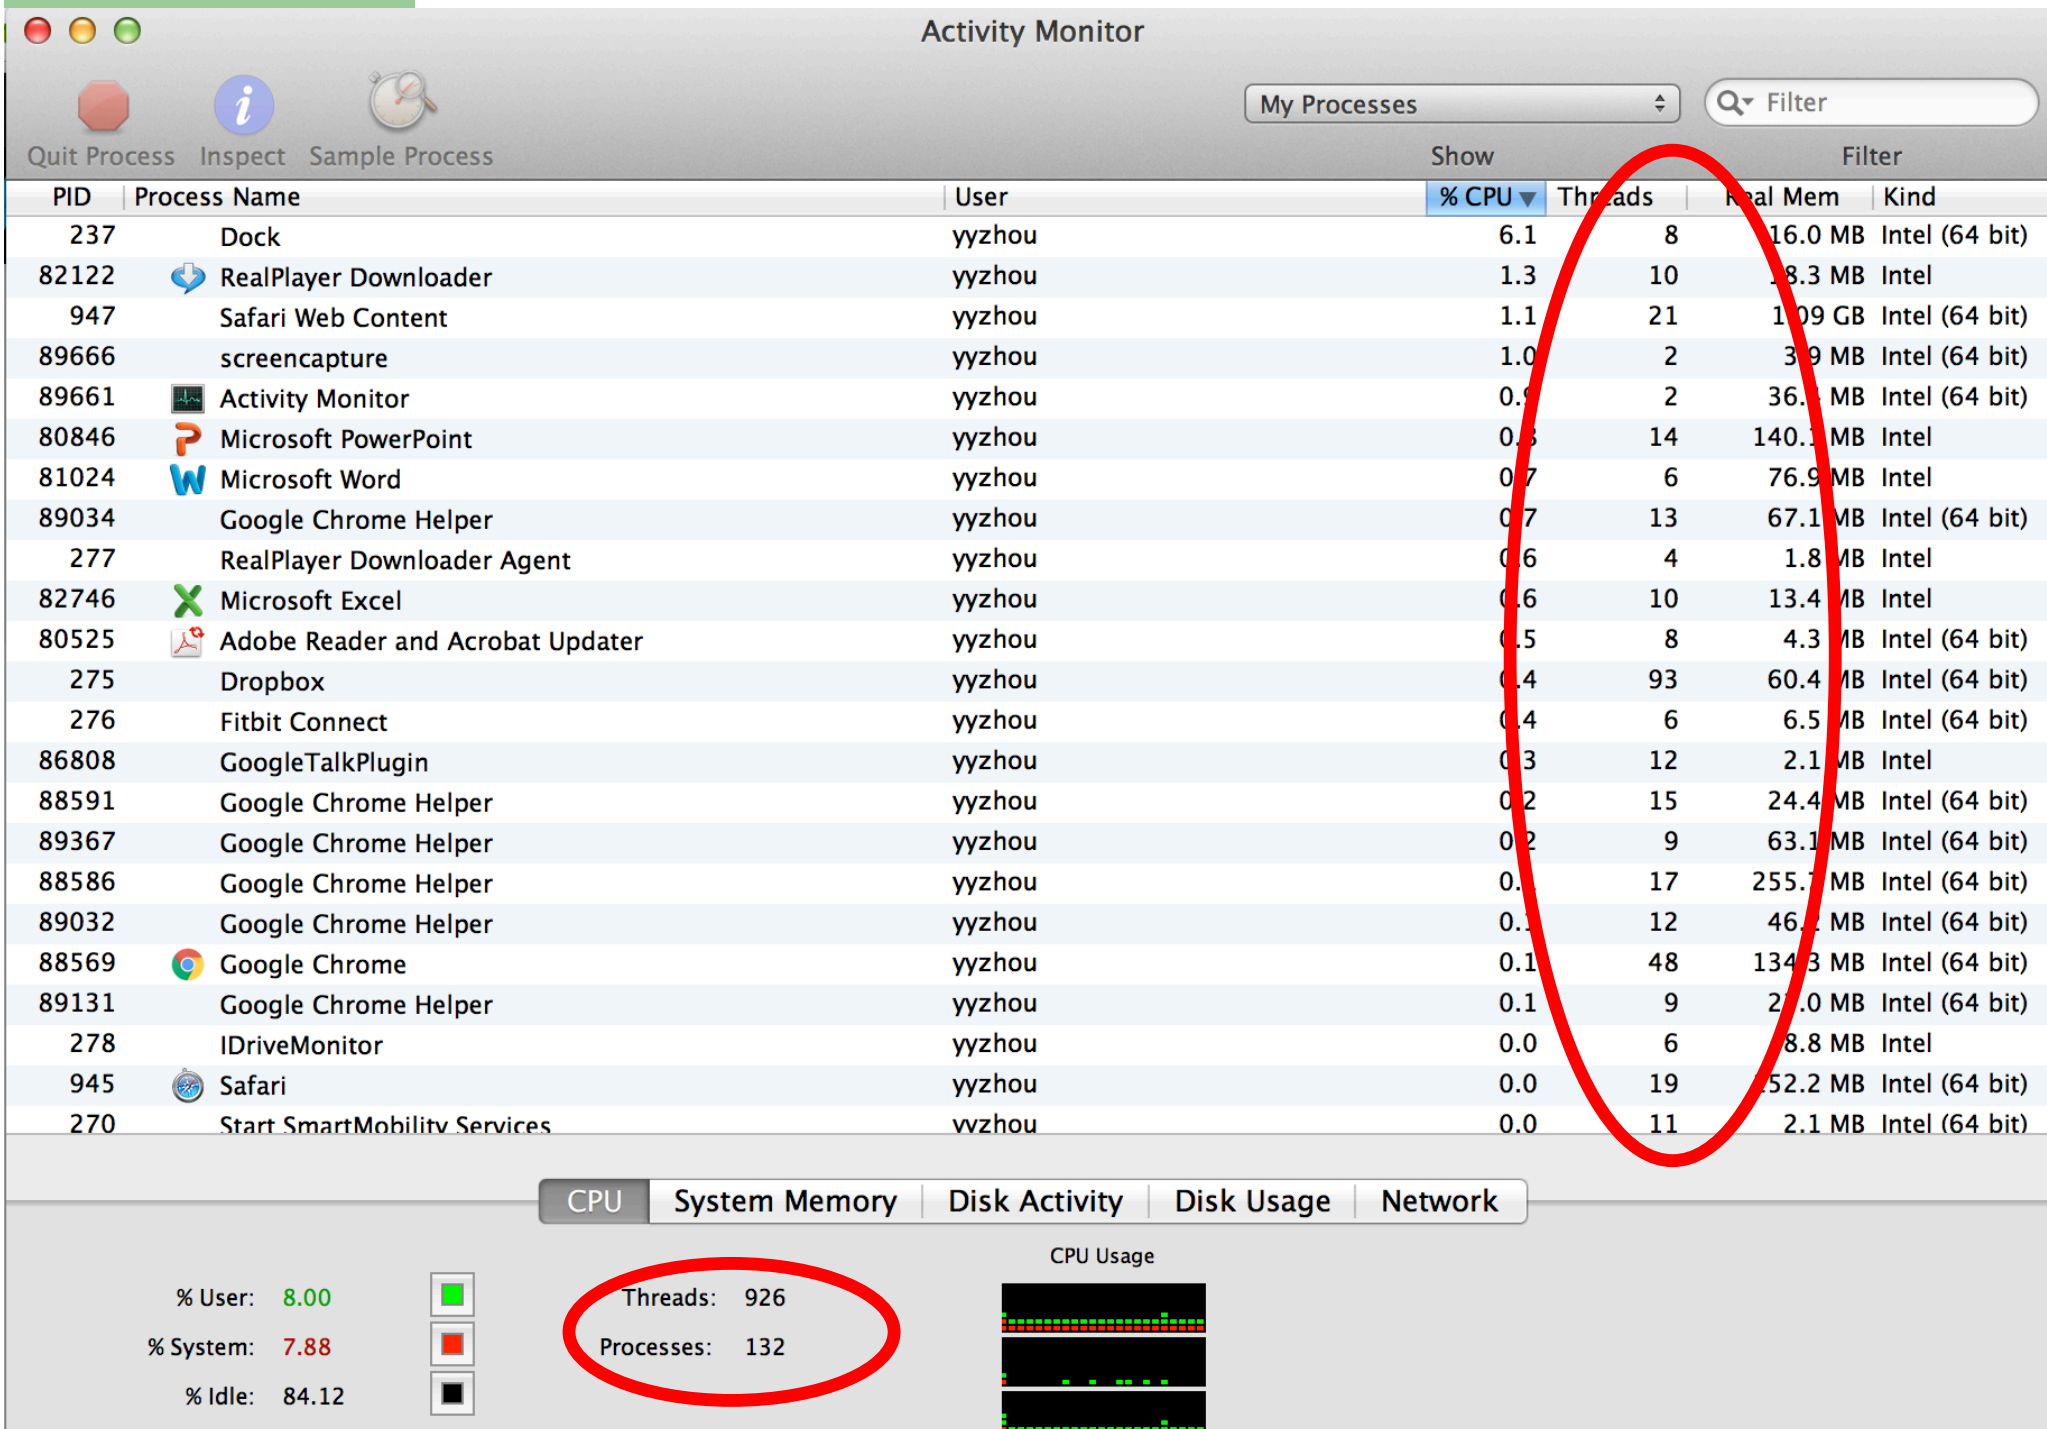
Windows Thread Lists from Performance Monitor



Windows Performance Analyzer



Mac OS – Activity Monitor



Thread Information on Linux

- Process information:
 - Read **/proc/[your PID]/stat** file
- Thread information (2.6 kernel):
 - Read **/proc/[your PID]/task/[thread ID]/stat**

Kernel-managed Threads

- We have taken the execution aspect of a process and separated it out into threads
 - To make concurrency cheaper
- As such, the **OS** now **manages** threads *and* processes
 - All thread operations are implemented in the kernel
 - The OS schedules all of the threads in the system
- OS-managed threads are called **kernel-level threads** or **Kernel managed threads** or **lightweight processes**
 - NT: threads
 - Solaris: lightweight processes (LWP)
 - POSIX Threads (pthreads): `PTHREAD_SCOPE_SYSTEM`

Kernel-managed Thread Limitations

- Kernel-managed threads make concurrency much cheaper than processes
 - Much less state to allocate and initialize
- However, for fine-grained concurrency, kernel-managed threads still suffer from too much overhead
 - Thread operations still require system calls
 - Ideally, want thread operations to be **as fast as a procedure call**
 - Kernel-level threads have to be general to support the needs of all programmers, languages, runtimes, etc.
- For more fine-grained concurrency, need even “cheaper” threads

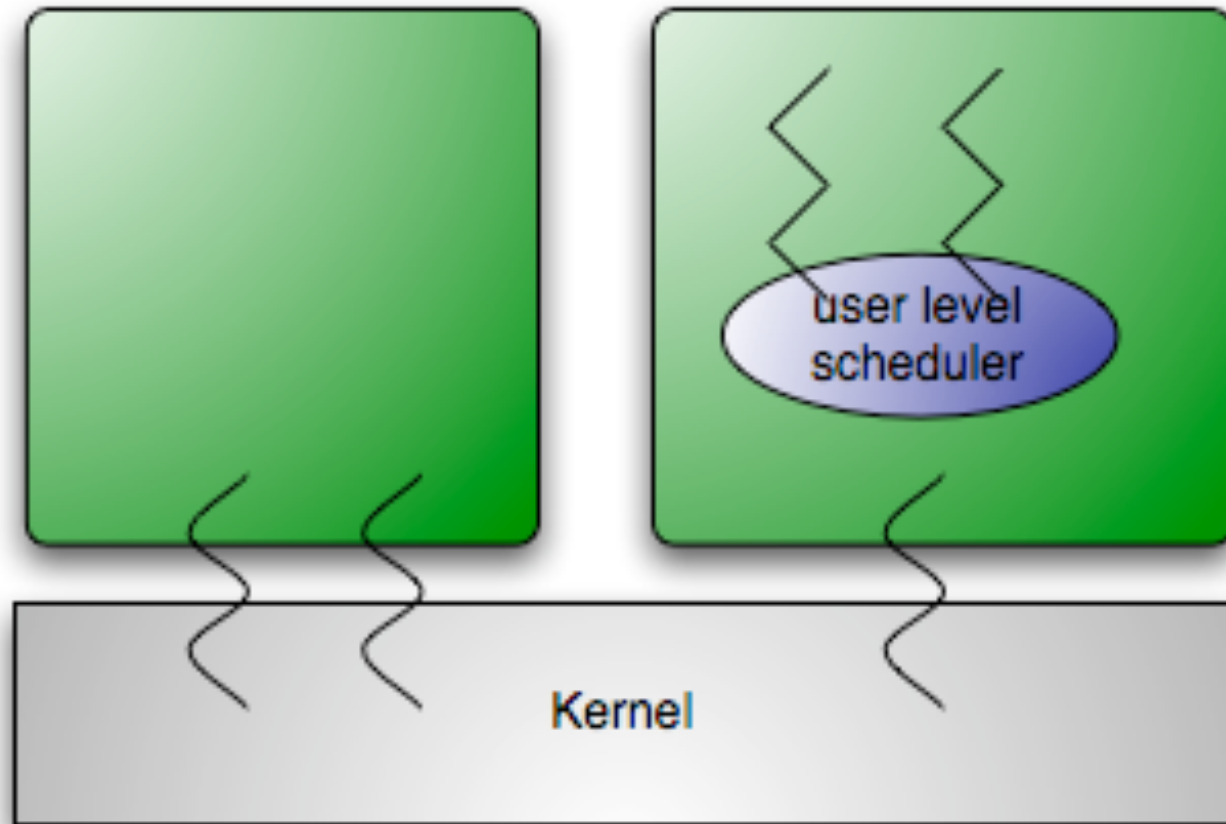
User-Level-Managed Threads

- To make threads cheap and fast, they need to be implemented at user level
 - Kernel-level managed threads are managed by the OS
 - User-level managed threads are managed entirely by the **run-time system (user-level library)**
- User-level-managed threads are small and fast
 - A thread is simply represented by a PC, registers, stack, and small **thread control block** (TCB)
 - Creating a new thread, switching between threads, and synchronizing threads are done via **procedure call**
 - No kernel involvement
 - User-level managed thread operations **100x faster** than kernel managed threads
 - pthreads: **PTHREAD_SCOPE_PROCESS**

User level threads

Kernel level Managed threads

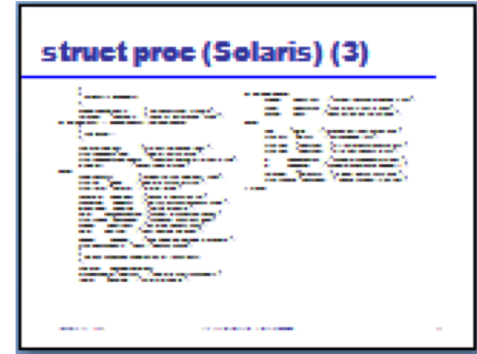
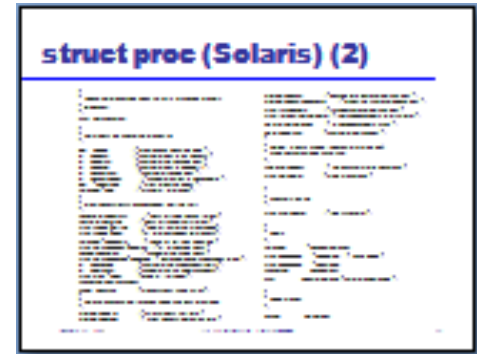
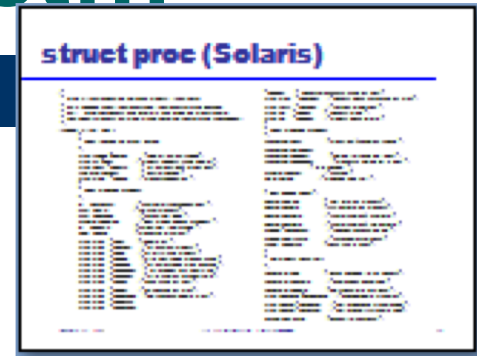
User level Managed threads



Small and Fast...

- Nachos thread control block

```
class Thread {  
    int *stack;  
    int *stackTop;  
    int machineState[MachineStateSize];  
    ThreadStatus status;  
    char *name;  
    <Methods>  
};
```



User Level Managed Thread Limitations

- But, user-level managed threads are not a perfect solution
 - As with everything else, they are a tradeoff
- User-level managed threads are **invisible** to the OS
 - They are not well integrated with the OS
- As a result, the OS can make poor decisions
 - Scheduling a process with idle threads
 - Blocking a process whose thread initiated an I/O, even though the process has other threads that can execute
 - Unscheduling a process with a thread holding a lock
- Solving this requires communication between the kernel and the user-level thread manager

Tradeoffs between the two

- Kernel-level managed threads
 - Integrated with OS (informed scheduling)
 - Slow to create, manipulate, synchronize
- User-level managed threads
 - Fast to create, manipulate, synchronize
 - Not integrated with OS (uninformed scheduling)
- Understanding the differences between kernel and user-level managed threads is important
 - For programming (correctness, performance)

Kernel and User Threads

- Or use **both** kernel and user-level threads
 - Can associate a user-level thread with a kernel-level thread
 - Or, multiplex user-level threads on top of kernel-level threads
- Java Virtual Machine (JVM) (also pthreads)
 - Java threads are user-level threads
 - On older Unix, only one “kernel thread” per process
 - Multiplex all Java threads on this one kernel thread
 - On NT, modern Unix
 - Can multiplex Java threads on multiple kernel threads
 - Can have more Java threads than kernel threads

Implementing Threads

- Implementing threads has a number of issues
 - Interface
 - Context switch
 - Preemptive vs. non-preemptive
 - Scheduling
 - Synchronization (next lecture)
- Focus on user-level managed threads
 - Kernel-level managed threads are similar to original process management and implementation in the OS
 - What you will be dealing with in Nachos
 - *Not only will you be using threads in Nachos, you will be implementing more thread functionality*

Sample Thread Interface

- `thread_fork(procedure_t)`
 - Create a new thread of control
 - Also `thread_create()`, `thread_setstate()`
- `thread_stop()`
 - Stop the calling thread; also `thread_block`
- `thread_start(thread_t)`
 - Start the given thread
- `thread_yield()`
 - Voluntarily give up the processor
- `thread_exit()`
 - Terminate the calling thread; also `thread_destroy`
- `Thread_join(t)` or `t.join()`
 - causes the current thread to pause execution until t's thread terminates

Thread Scheduling

- The thread scheduler determines when a thread runs
- It uses queues to keep track of what threads are doing
 - Just like the OS and processes
 - But it is implemented at user-level in a library
- Run queue: Threads currently running (usually one)
- Ready queue: Threads ready to run
- Are there wait queues?
 - How would you implement `thread_sleep(time)`?

Non-Preemptive Scheduling

- Threads **voluntarily** give up the CPU with `thread_yield`

Ping Thread

```
while (1) {  
    printf("ping\n");  
    thread_yield();  
}
```

Pong Thread

```
while (1) {  
    printf("pong\n");  
    thread_yield();  
}
```

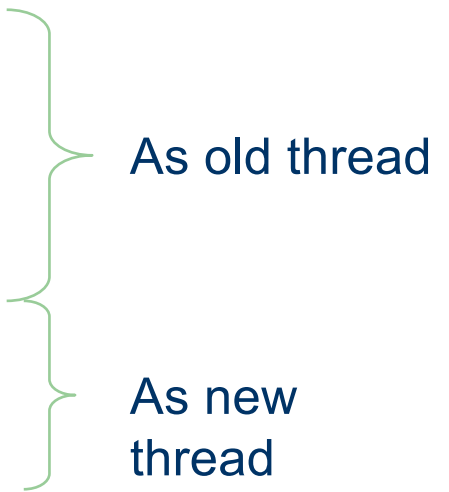
- What is the output of running these two threads?

thread_yield()

- Wait a second. How does thread_yield() work?
- The semantics of thread_yield are that it gives up the CPU to another thread
 - In other words, it **context switches** to another thread
- So what does it mean for thread_yield to return?
 - It means that *another thread* called thread_yield!
- Execution trace of ping/pong
 - `printf("ping\n");`
 - `thread_yield();`
 - `printf("pong\n");`
 - `thread_yield();`
 - ...

Implementing thread_yield()

```
thread_yield() {  
    thread_t old_thread = current_thread;  
    current_thread = get_next_thread();  
    append_to_queue(ready_queue, old_thread);  
    context_switch(old_thread, current_thread);  
    return;  
}
```



As old thread

As new thread

- The magic step is invoking context_switch()
- Why do we need to call append_to_queue()?

Thread Context Switch

- The context switch routine does all of the magic
 - Saves context of the currently running thread (`old_thread`)
 - Push all machine state onto its stack (*not* its TCB)
 - Restores context of the next thread
 - Pop all machine state from the next thread's stack
 - The next thread becomes the current thread
 - Return to the NEW thread
- This is all done in assembly language
 - It works **at** the level of the procedure calling convention, so it cannot be implemented using procedure calls
 - See [code/threads/switch.s](#) in Nachos

Preemptive Scheduling

- Non-preemptive threads have to voluntarily give up CPU
 - A long-running thread will take over the CPU
 - Only voluntary calls to `thread_yield()`, `thread_stop()`, or `thread_exit()` causes a context switch
- **Preemptive scheduling** causes an **involuntary** context switch
 - Need to regain control of processor asynchronously
 - How?
 - Use timer/alarm interrupt

Blocking Vs. non-blocking System Calls

- Blocking system call
 - Usually I/O related: read(), fread(), getc(), write()
 - Doesn't return until the call completes
 - The process/thread is switched to blocked state
 - When the I/O completes, the process/thread becomes ready
 - Simple
 - Real life example: attending a lecture
- Using non-blocking system call for I/O
 - Asynchronous I/O
 - Complicated
 - The call returns once the I/O is initiated, and the caller continue
 - Once the I/O completes, an interrupt is delivered to the caller
 - Real life example: apply for job

Threads Summary

- The operating system as a large multithreaded program
 - Each process executes as a thread within the OS
- Multithreading is also very useful for applications
 - Efficient multithreading requires fast primitives
 - Processes are too heavyweight
- Solution is to separate threads from processes
 - Kernel-level managed threads much better, but still significant overhead
 - User-level managed threads even better, but not well integrated with OS
- Now, how do we get our threads to correctly cooperate with each other?
 - Synchronization...